



# A method for machine learning generation of realistic synthetic datasets for validating healthcare applications

Health Informatics Journal

2022, Vol. 0(0) 1–16

© The Author(s) 2022

Article reuse guidelines:

[sagepub.com/journals-permissions](https://sagepub.com/journals-permissions)

DOI: 10.1177/14604582221077000

[journals.sagepub.com/home/jhi](https://journals.sagepub.com/home/jhi)**Theodoros N Arvanitis** 

Institute of Digital Healthcare, WMG, University of Warwick, Coventry, UK

**Sean White**

Clinical Assurance Team, NHS Digital, Leeds, UK

**Stuart Harrison**

Institute of Digital Healthcare, WMG, University of Warwick, Coventry, UK

**Rupert Chaplin**

Data Science and Innovation, NHS Digital, London, UK

**George Despotou** 

Institute of Digital Healthcare, WMG, University of Warwick, Coventry, UK

## Abstract

Digital health applications can improve quality and effectiveness of healthcare, by offering a number of new tools to users, which are often considered a medical device. Assuring their safe operation requires, amongst others, clinical validation, needing large datasets to test them in realistic clinical scenarios. Access to datasets is challenging, due to patient privacy concerns. Development of synthetic datasets is seen as a potential alternative. The objective of the paper is the development of a method for the generation of realistic synthetic datasets, statistically equivalent to real clinical datasets, and demonstrate that the Generative Adversarial Network (GAN) based approach is fit for purpose. A generative adversarial network was implemented and trained, in a series of six experiments, using numerical and categorical variables, including ICD-9 and laboratory codes, from

## Corresponding author:

George Despotou, Institute of Digital Healthcare, WMG, University of Warwick, Coventry CV4 7AL, UK.

Email: [g.despotou@warwick.ac.uk](mailto:g.despotou@warwick.ac.uk)

Creative Commons CC BY: This article is distributed under the terms of the Creative

Commons Attribution 4.0 License (<https://creativecommons.org/licenses/by/4.0/>) which

permits any use, reproduction and distribution of the work without further permission provided the

original work is attributed as specified on the SAGE and Open Access pages (<https://us.sagepub.com/en-us/nam/open-access-at-sage>).

three clinically relevant datasets. A number of contextual steps provided the success criteria for the synthetic dataset. A synthetic dataset that exhibits very similar statistical characteristics with the real dataset was generated. Pairwise association of variables is very similar. A high degree of Jaccard similarity and a successful K-S test further support this. The proof of concept of generating realistic synthetic datasets was successful, with the approach showing promise for further work.

## Keywords

Generative adversarial networks, certification, machine learning, realistic synthetic datasets, safety

## Introduction

Digital health has seen a continuously increasing number of innovative applications, aiming to improve all aspects of one's health and care; such as safety, efficacy and monitoring of care plans, empowerment of the patient to manage their own condition, as well as discovery of new clinical knowledge. These innovations have been made possible by applying state-of-the-art computer science, data science and software engineering technologies in healthcare, resulting in applications such as automated diagnosis, self-monitoring, telehealth and clinical decision support. Applications can vary from simple statistics viewers to symptom checkers and diagnostic services, used in both primary and secondary care. Their integration with clinical pathways can be seen as introducing Clinical Decision Support (CDS). Failures in their operation may cause harm to patients. For example, by offering incorrect advice, by making the wrong diagnosis or recommendation to a healthcare professional. This potential for harm is increasingly recognised and has been incorporated in regulation, where software is seen as a medical device.<sup>1-3</sup> Software applications may expose patients to risks, due to unintended or erroneous behaviour, or lack of clinical validation that may create unknowns, in whether these applications are fit-for-purpose and acceptably effective. Understanding risks and establishing clinical validation is necessary, for software manufacturers to be able to meet certification and regulation requirements, in order to offer their products to the healthcare system and patients. Clinical validation entails comparing the application against datasets, typical of the patients who will use it. This can be problematic as large patient datasets are not readily available to manufacturers due to regulations and law governing patient privacy.<sup>4</sup> The problem is aggravated, as datasets need to be suitable for the context of an application, with manufacturers often resorting to buying compiled data. An increasing number of applications requiring datasets also increases demand and hence waiting times. Furthermore, even if manufacturers can access data, either through a source or produced by themselves, validation is still difficult as the regulator will need to have access to the same dataset (or a common dataset based on which to exchange information about the fitness of the application). As a result, patients may be deprived of digitally enabled applications improving the healthcare quality they receive. Use of Realistic Synthetic Datasets (RSDs) is seen as a promising solution,<sup>5-9</sup> addressing privacy challenges, whilst overcoming issues with alternatives; such as anonymized data that may skew the results of validation due to missing fields.<sup>4</sup> Natural Language Processing can enhance datasets by adding synthetic clinical notes, based on real records.<sup>10,11</sup> RSDs consist of software-generated data points, which overall demonstrate equivalent statistical properties as a real clinical dataset. In the context of clinical validation, use of the two datasets should result in the same conclusions, with the same confidence. Contrary to using de-identified or anonymized datasets, RSDs: (1) do not need prolonged preparation and approval process, something that is required even with anonymized data;

(2) can include variables that may be considered sensitive with respect to patients' privacy and are not included in anonymized and de-identified datasets; (3) are highly resistant to cross-referencing with other datasets (although some concerns still remain to be addressed). The most common approach, to developing synthetic datasets, looks at the associations and the distributions of variables of interest, and develops probabilistic models that will then generate the synthetic data.<sup>12-15</sup> A potential drawback of this is that they often use a driving variable, which is generated first, and then drives the associations with the others,<sup>8</sup> potentially skewing the dataset to conditions represented by these variables, and omitting other variables that may hide unearthed associations.

Machine Learning has been prominent in producing large RSDs, with similar statistical qualities to a real dataset on which they are trained.<sup>16-18</sup> Generative Adversarial Networks are a ML approach based on neural networks, recently recognized as able to accurately mimic real datasets.<sup>19-22</sup> GANs, in general, will start their training without relying on statistical models representing real dataset, and will establish the statistical properties of the real dataset.<sup>23</sup> This allows GANs to be 'dataset agnostic' and transferable across multiple datasets. Evaluating the performance of GANs requires focussing on statistical properties of the real and the synthetic dataset.<sup>24,25</sup> This paper presents the results of a Realistic Synthetic Dataset Generation Method (RSDGM) using GANs, which generated a synthetic dataset suitable for validation of digital health applications.

## Objectives

A proof of concept of producing realistic statistically equivalent (to real data), large-scale, scalable, machine-learning generated datasets, for validation of healthcare applications used across all levels:

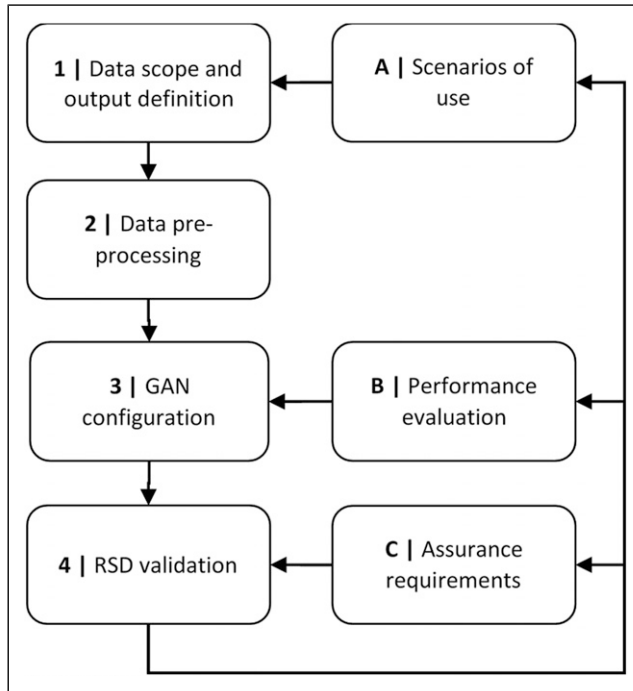
- *Realistic*: the dataset will need to be statistically equivalent to real dataset.
- *Large-scale*: the dataset will need to consist of a large number of entries and variables.
- *Scalable*: the method should be scalable for larger datasets.
- *Machine learning generated*: machine-learning (GANs) should be used to generate the dataset, and the optimum hyper-parameters need to be identified.
- *Validation of healthcare applications*: the method should generate evidence justifying the suitability of the RSD to validate healthcare applications.
- *Levels of care*: RSDs for primary as well as secondary care.

The scope of this work focused on examining the feasibility, and effectiveness of the techniques, and stopped short of packaging and deploying the data for mass use by multiple applications.

## Method

The method (Figure 1) consists of four (4) main steps (1, 2, 3 and 4) and three (3) contextual steps (a, b and c). The contextual steps provided the necessary framework for making methodological decisions.

The method was defined iteratively over six main experiments (Table 1), each of which consisted of numerous runs, which allowed testing aspects such as the GAN hyper-parameters. For example, changing the hidden layers twice in experiment #3, which tested three different GANs in parallel, would result in six runs. The project did not need ethics approval; all datasets used were freely and publicly available (see step 1 description). All experiments were run on a computer with 2x Intel Xeon Gold 6144 3.5GHz, 3x Nvidia Quadro RTX5000, 8x16GB DDR4 2666MHz RDIMM ECC and Windows 10 Professional for workstations. Implementation of each experiment was done in the



**Figure 1.** Overview of the Realistic Synthetic Dataset Generation Method (RSDGM).

Anaconda, in Python 3.6 using Tensorflow 2. Statistical validation was done in the same framework using the *scipy*, *math*, *matplotlib*, *pandas*, *seaborn* and *numpy* packages. Step 1 explains the implementations used in each experiment.

### *Data Scope and Output Definition*

This step defined the content and purpose of the dataset, providing the business goals for the experiments, and specifies the success requirements of the RSD:

- The method needs to be easily scalable to re-train the dataset with additional variables.
- Allow modelling of large number of variables.
- Establish associations amongst all variables.
- Produce both numerical and categorical values.
- Allow the generation of clinical codes (e.g. ICD-9).

Table 1 shows an overview of the experiments performed as part of developing the RSDGM. Three different datasets (NHS HES,<sup>26</sup> NHS A&E,<sup>27</sup> and MIMIC III<sup>28</sup>) were used to ensure that the method produced the same results in different datasets, which was a project requirement. MIMIC III was considered the most complete dataset, and being open, was preferred for the final experiments.

Experiment #1 focused on a vanilla GAN, with numerical values. Experiments #2, #3, #4 compared the implementations of three different GANs; a vanilla GAN (GAN),<sup>29</sup> a conditional GAN (CGAN),<sup>30</sup> and a Wasserstein GAN (WGAN),<sup>31</sup> subsequently implementing gradient penalty

**Table I.** Overview of the six experiments; the dataset used, the type of GAN implemented, the complexity added in each experiment, sample size, number of variables (N: numeric, BC: Binary Categorical, C: categorical; brackets show the number of minimum and maximum size of categorical values) and the approximate duration of each run in the experiment.

Exp	Dataset	GAN	Size	Variables	Additional Complexity	Appx. Duration
1	HES	GAN	150	10xN	Simple GAN, Numerical values.	10 min
2	MIMIC-III	GAN, CGAN, WGAN	130	2xN, 6xBC	Evaluate multiple GAN implementations. Increase in epochs to understand progression and performance.	25 min
3	MIMIC-III	GAN, CGAN, WGANP	5000	3xN, 7xC(2,12)	Multiple GAN implementations. Bigger sample to test performance. Training of categorical values.	45 min
4	A&E	GAN, CGAN, WGANP	1000	3xN, 7xC(2,4)	Multi-dataset validation check.	20 min
5	MIMIC-III	WGANP	150	1xC(-,513)	Implementation of ICD-9 codes. High number of values of categorical variables.	8 h
6	MIMIC-III	WGANP	379	2xC(-,1357)	Implementation of lab test codes, and ICD-9 based admission.	12 h

(WGANP).<sup>32</sup> The WGANP was considered to have the best performance, and was selected for the last experiments, but with the added complexity of generating ICD-9 codes (exp. #5), and lab test codes (exp. #6).

*Context A: Scenarios of use:* The main scenario was the validation of clinical decision support applications, providing recommendations to patients or healthcare professionals.<sup>33</sup> The scope encompassed both primary and secondary care. Primary care datasets focus on a relatively low number of variables (e.g. weight, age, diagnosis) generally describing bigger population. Secondary care datasets need a larger number of variables, relating to the protocol followed during the encounter of the patient, including lab tests.

### Data pre-processing

Pre-processing removed missing and duplicate data. Multiple categorical features were interpreted as a collection of binary features, with one-hot encoding. The input to this transformer is a matrix of integers, denoting the values of the categorical features. The output is a sparse matrix, where each column corresponds to one value of a feature. All variables of the training data are rescaled by applying standardisation. We multiply the forecast value of the standardised input by the standard deviation calculated in the original series, and then add the mean. Label encoder was used for the categorical variables to generate their combined pairwise correlations.

There are two main challenges in incorporating lab codes and ICD-9 variables in the generation method. The first is the complexity of the variables. There are thousands of ICD9 codes in the

In the experiments, we generated ICD-9 codes per hospital admission, using the `DIAGNOSIS_ICD.csv` file of MIMIC-III. In this, each patient has multiple hospital stays and each stay is associated with a unique ICD-9 code. The diagnosis variable was expanded in a separate matrix with length equal to the number of unique codes, containing binary values capturing the presence or not of an ICD-9 code (Table 2). The matrix only captured the presence of a diagnosis without adding other information, such as weight of primary diagnosis and co-morbidities, which would require review by physicians.

## GAN design and configuration

The Adam optimiser<sup>34</sup> showed best performance in early evaluation, when compared to RMSProbOptimiser as part of (a non-exhaustive) comparison between implementation approaches in literature. In the last experiments, two hidden layers, both for the discriminator (D) and the generator (G), were used. The number of neurons was the *sum of input and output layers multiplied by 2/3 of the value given between the input and output sizes*. Mean squared loss for numerical variables (e.g. age) and cross entropy loss for categorical variables were initially used. However, Wasserstein loss was later used, improving the results. The first experiment used the Leaky ReLU activation function, whereas subsequent experiments used Rectified Linear Units (RELU), as the activation function of the Generator, except of the output layer where *tanh* was used for numerical variables. For the discriminator, ReLU was used for all activation functions, except for the output layer, where the sigmoid function for categorical and *tanh* for numeric features were used.<sup>35</sup> The number of epochs ranged from 10,000 to 50,000. A trial implementation tested the network up to 300,000 epochs, showing that, after 50,000 epochs changes were insignificant. A target of 50,000

**Table 2.** Excerpt of patient – ICD9 code matrix.

Subject ID	Hadm ID	ICD-9 Codes									
		042	135	486	1917	30401	40301	E8502	V502	...	V721
78	100536	1	0	0	0	1	0	1	0	...	0
41	101757	0	0	0	1	0	0	0	0	...	0
109	102024	0	0	0	0	0	1	0	0	...	0
109	172335	0	0	1	0	0	1	0	0	...	0
...	...	...	...	...	...	...	...	...	...	...	...
16	103251	0	0	0	0	0	0	0	1	...	1

**Table 3.** GAN configurations for each experiment (D and G refer to the discriminative and generative networks, respectively).

E#	Epochs	Learning rate	Hidden layers	Neurons	Batch size	Optimiser
1	10,000	$1 \times 10^{-3}$	3 D, 2G	D[8, 8, 8, 1] G[8, 8]	150	RMSProbOptimiser
2	30,000	$1 \times 10^{-4}$	1 D, 1G	D[18,1] G[21]	130	Adam
3	50,000	$1 \times 10^{-4}$	1 D, 1G	D[18,1] G[21]	5000	Adam
4	10,000	$1 \times 10^{-4}$	1 D, 1G	D[18,1] G[21]	1000	Adam
5	30,000	$1 \times 10^{-4}$	2 D, 2G	D[256,128,1] G[256, 128]	150*	Adam
6	50,000	$1 \times 10^{-4}$	2 D, 2G	D[256,128,1] G[256, 128]	379*	Adam

epochs was considered a good balance between training and computing power. The number of discriminator iterations per generator iteration is 10, and the Gradient Penalty (GP) (*lambda-penalty coefficient*) is 10. The raw generated data values were continuous in range 0–1, converted to binary (0 or 1) through rounding for categorical data classification. For numeric features, the output of generator is de-standardised, to make it amenable for manual review, as a sanity check of the output of the GAN.

*Context B: Performance evaluation:* As this was an exploratory, proof-of-concept study, the method evolved along with the results of the experiments. The performance evaluation, used for this, was a meta-process. It analysed the results and shaped the final form of the experiments in Table 3. After each experiment, this approach identified parameters for rapid trial and error, as well as the improvements in subsequent experiments (e.g. increase of epochs).

### Synthetic dataset validation

This step examined whether the RSD would be fit for purpose for the objective of the study and the scenario identified in context A. From the early stages of the experiment, it was clear that the validation of the synthetic dataset is a major challenge. Although the dataset contains different data points, its overall qualities needed to be equivalent to the real dataset. Qualifying the justification for equivalence requires understanding the validation context, in order to identify suitable evidence. This was achieved by a parallel, assurance process (Context C), using a justification outline,<sup>36</sup> identifying the evidence needed to be generated to support that justification. The aspects needed support by evidence were recognised by conducting a safety assessment, identifying potential risks using the RSDGM. The identified justification comprises of three main arguments. Firstly evidence needs to demonstrate that the RSD is a high fidelity representation of clinical knowledge (e.g. prevalence of conditions), including potential relationships not represented in existing knowledge. Secondly, there needs to be evidence supporting that the synthetic and the real datasets exhibit almost identical statistical properties (e.g. associations amongst variables). The final argument of the justification focuses on the technical correctness and appropriate application of the generation method. The first argument is beyond the scope of this paper, as it focuses on the use of the RSD. The second and third arguments resulted in a process identifying evidence, which would convincingly justify the validation of the RSD. The loss function, as well as scatter plots of the datasets through regular intervals of training the GAN, demonstrates that the network performed as predicted by theory. Visual comparison of the data entries, association tables, jaccard similarity (using sklearn.metrics), and a K-S test (using scipy.stats ks\_2samp to compare the generated and original dataset), were identified as convincing evidence to support statistical equivalence.

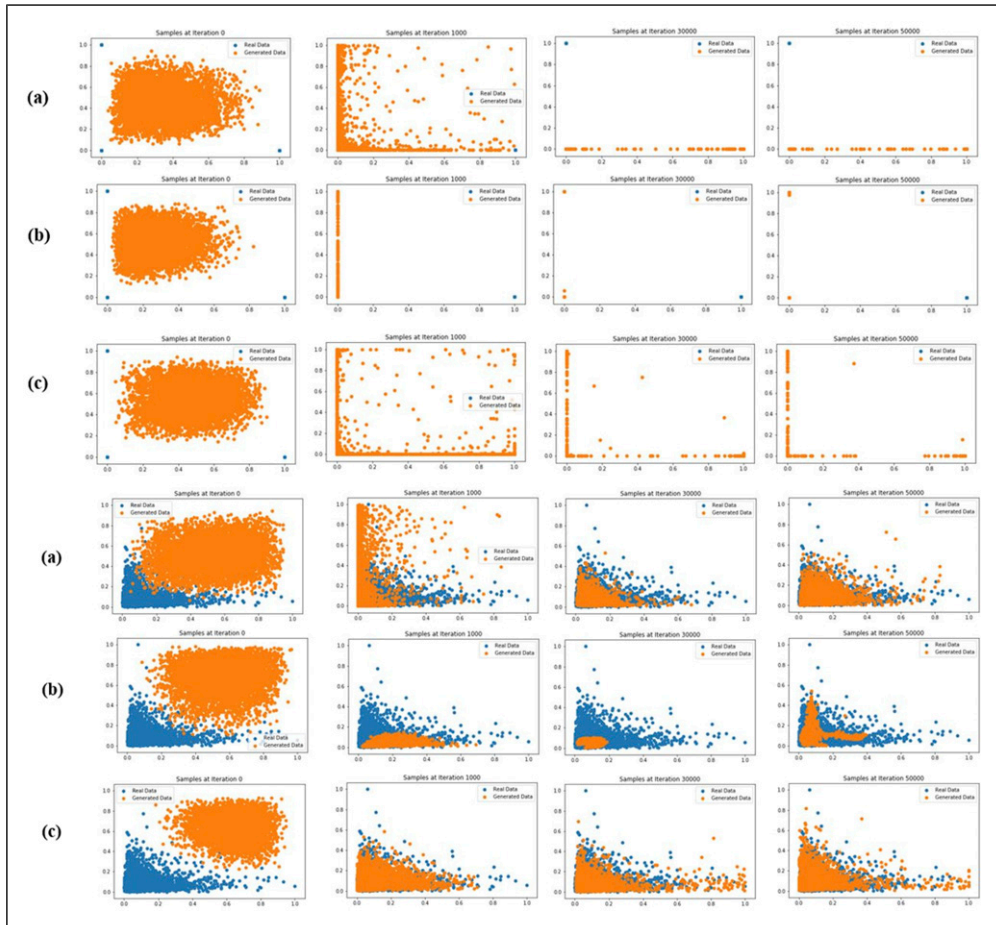


## Results

### Comparison of the GAN implementations

Wasserstein loss consistently produced the best results in the experiments, where multiple GANs were implemented. Figure 2 presents the results of selected features (top: GENDER & ETHNICITY, bottom: EDOUT\_EDREG\_TIME & OUTIME\_INTIME\_WARDS) from experiment #3 (experiment #3 included 10 features from the ADMISSION, PATIENT and ICUSTAYS tables of the MIMIC-III dataset; 7 were categorical and 3 numeric, trained for 50,000 epochs). The scatter plots show the real data (points in blue) overlapped with the generated data (points in orange).

The WGAN performed best with the two types of variables. The vanilla GAN did not manage to perform well generating categorical values completely missing one dimension, whereas the CGAN



**Figure 2.** Scatter plots showing the convergence of the synthetic (orange) dataset with the real (blue) dataset as the number of epochs increases. (bottom) – Two selected numerical features for 50,000 epochs: (a) GAN, (b) CGAN and (c) WGANGP; (top) – Two selected categorical features for 50,000 epochs: (a) GAN, (b) CGAN and (c) WGANGP.



was very poor generating numerical values. Figure 3 presents the losses functions (10 iteration intervals; Generator in blue and Discriminator in orange). The WGAN implementation gave a better loss function plot with the Generator achieving good response, and the Discriminator not able to confidently discern true and false data.


### RSDGM validation framework

The validation approach of the RSDGM consists of (a) loss analysis, (b) visualisation of real and generated data, (c) correlations comparison and (d) similarity measurement. The loss functions of the discriminator and generator were checked for two common problems: (a) the generator collapsing and generating only one sample, and (b) the generator simply memorizing or being too similar to the training data. The discriminator loss is trying to minimize the classification error of the discriminator and generator loss is trying to maximize the classification error of the discriminator. Although the loss function is not providing direct evidence about the statistical equivalence of the two datasets, it was particularly useful in early experiments, as it offered confidence that the network was operating as intended. It offered some early evidence and confidence about the algorithmic determinism of the implementation.

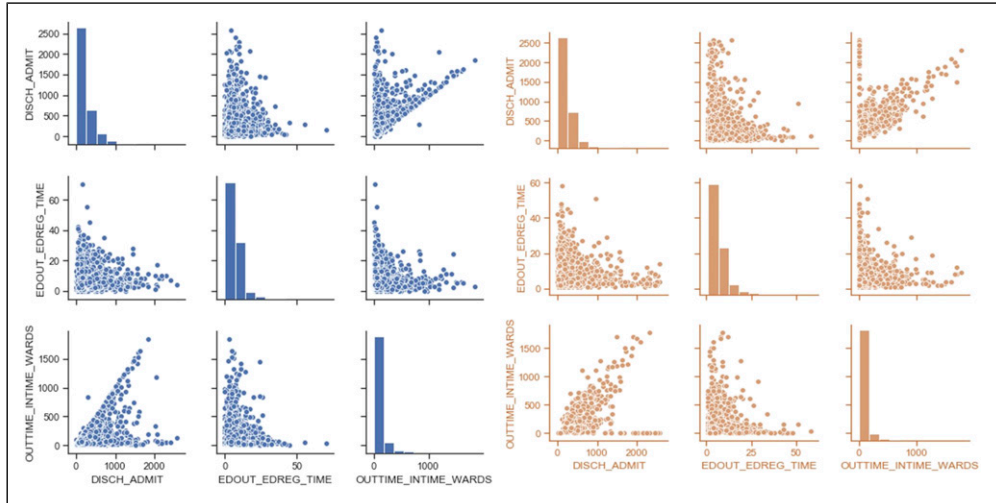
Visualisation of data was the first type of evidence produced, able to demonstrate statistical equivalence. It offers an intuitive way to understand the datasets, particularly effective in smaller size experiments, as it gave sufficient confidence about the similarity of the datasets; whilst giving the opportunity to spot potential outliers or over/under fitting. Furthermore, the plots gave an understanding of the behaviour of the network during its training, being plotted every 1000 iterations. They visualise how the Generator network starts, with a random initial mapping between the input and dataset vector space, and then gradually evolves to resemble the real dataset.

A correlation matrix of each dataset was computed in each experiment. The correlation matrix was visualized as a heat map. Correlation matrices provide confidence that the synthetic dataset has maintained an equivalent association amongst variables. Pearson correlation was applied for count or label-encoded categorical features. Spearman correlation was applied for binary features, as well as for experiments with mixed variables.

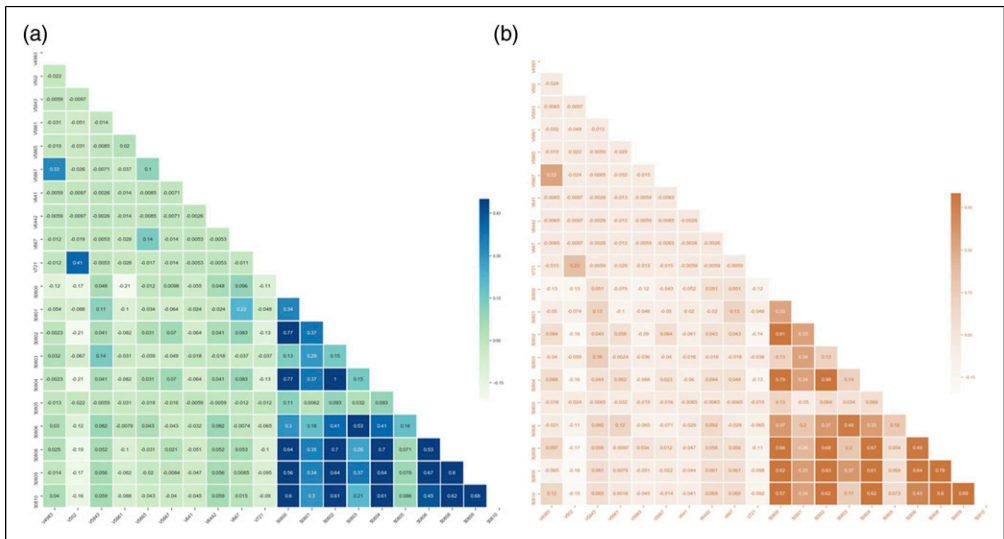
Ensuring whether the RSD learned the distribution of each dimension acceptably, the Mann–Whitney  $U$  test was applied. The test was used to compare whether the distributions of each variable of real and generated dataset come from the same population. Furthermore, two-sample Kolmogorov-Smirnov (K-S) test, which is a non-parametric test that compares the cumulative distributions of two data sets, was used to compare real and generated datasets. The null hypothesis of this test is that both samples originate from a population with the same distribution.



**Figure 3.** Generator and Discriminator losses for 50,000 epochs: (a) GAN, (b) CGAN and (c) WGAN. WGAN resulted in a more meaningful loss, as the synthetic data converged with the real data, reducing its loss, and the discriminator not being able to confidently classify data.




**Figure 4.** Comparison of the distribution of a selection of numerical values in the two datasets (blue: real, orange: synthetic).



**Figure 5.** Correlation matrices of (a) real and (b) generated data for 10 ICD9-codes and 10 Lab item codes, shows a very similar but not identical associations amongst features in the two datasets.

Finally, Jaccard similarity indices were used to compare associations, limited to absence/presence data. It is a measure of similarity for the two sets of data, with a range from 0% to 100%. The higher the percentage, the more similar the two populations. We considered dichotomic variables with 0 or 1 values, absence or presence of ICD9-codes or Lab item code per patient admission, and calculated the similarities between real and generated data for each code.



**Figure 6.** (left) Histogram of Jaccard scores for lab item codes, ICD-9 codes and total number of codes; (right) Scatter plot of dimension-wise probability results of real binary data (x-axis) vs. synthetic counterpart (y-axis) produced.

### WPGAN generated realistic synthetic dataset

The RSDGM used the MIMIC III dataset<sup>28</sup> as the real dataset to train on. The experiment (#6) that resulted in the final generated dataset was run using the following hyper-parameters: (a) 50,000 epochs, (b) a ReLU activation function, (c) learning rate of  $10^{-4}$ , (d) 2 hidden layers D [256, 128, 1] G [256, 128], (e) Wasserstein distance as loss function, (f) the Adam optimized and (g) penalty gradient 10. The experiment used the DIAGNOSIS\_ICD and LABEVENTS tables of the MIMIC-III dataset. In total, 1357 codes (944 ICD9 + 413 Lab item) and 379 common unique hospital admissions of 300 patients were used as the real dataset.

Figure 4 illustrates the distribution of a sample of the variables in the RSD and the real dataset. The synthetic dataset shows a good representation of both numerical and categorical datasets with very similar distributions. The loss function (not illustrated) followed the expected response, similar to Figure 3 (right).

Figure 5 shows an extract of matrices of pairwise Spearman correlation of variables in the RSD and the real dataset. The matrices contain an extract of ICD-9 and lab codes. The very large dimensions of the full matrix do not allow complete graphical representation; nevertheless, further sample association matrices replicated the behaviour, also confirmed by a manual inspection of the associations. The RSD has preserved the associations between ICD-9 codes and the lab codes, of the original dataset.

Figure 6 (left) shows the Jaccard similarity of the ICD-9 and lab item codes. The majority of the variables indicated very high Jaccard similarity, and only a fewer lab codes resulted in low similarity. The K-S test failed to reject the null hypothesis with  $p$ -value ( $\alpha = 0.05$ ). Thus, there is no significant difference between the distributions for the two samples.

The scatter plot of dimension-wise probability results of real binary data versus synthetic counterpart is shown in Figure 6 (right). Experiment six, being the final experiment, provided evidence towards the proof of concept of the RSDGM, using a highly complex dataset, with complex associations.

## Discussion and conclusions

Realistic synthetic datasets are an approach recognised as promising, for validation and safety assurance of intelligent healthcare applications. This will overcome barriers of using datasets due to privacy concerns, enabling development of applications that may increase patient benefit. In our approach the machine learning model was trained using the real data, and then a synthetic dataset was generated that can be shared more widely to digital health innovators to validate their applications. This will allow developers to de-risk their applications early during their lifecycle, without the cost and delay of acquiring a real dataset. Exhaustive statistical analysis will be necessary to examine all aspects of equivalence, such as outliers; however, the statistical similarity of the synthetic and real datasets in this work, is considered to provide a significant degree of confidence to support innovation, until the more traditional and official stages of validation (and certification) of a digital health application (where a real dataset may be used for the highest degrees of confidence). Although the synthetic dataset was statistically equivalent to the real dataset, more work is needed to understand its role in the validation process. A risk assessment is necessary to understand potential issues, how they can be addressed, and the circumstances (e.g. early digital innovation re-risking) when potential uncertainty might be tolerated. Dataset-based validation of applications is a wider issue not limited to synthetic datasets; nevertheless in RSDs there is the additional steps of training and generating data that may introduce deficiencies in the dataset, such as outliers and representation of minority groups, that need more targeted testing. The work of this

project has identified a number of issues and tests as future work, as they lie beyond the scope of proof-of-concept work of this project. The GAN-based method, successfully generated a realistic synthetic dataset. Statistical tests demonstrated that the two datasets share very similar qualities. Some differences between the datasets were identified, particularly with respect to certain lab and ICD-9 codes. This was attributed to low frequency of certain conditions and lab tests. Bigger samples are needed to further explore this aspect. Although the datasets share very similar qualities, they are not completely identical. This was a positive finding as it meant that the GAN did not replicate the real seal dataset values, which would compromise privacy. Nevertheless, the degree of difference between the two datasets will need to be justified and accepted for the digital health application the RSD is intended to validate. Further validation from the point of view of clinical conditions would provide additional significant evidence on the equivalence of the two datasets. This would allow testing specific applications, and provide opportunities for expert (clinical) review of the dataset. Validation of applications using ML-generated RSDs heavily depends on contextual information about the application, as well as the generation of the dataset. A justification of use of the RSD is necessary, as it will allow the RSD developers to understand the safety implications of the generation process. Consequently, evidence need to be identified, supporting the justification. This will alleviate implications, offering evidence about fitness of the approach as a means of validation. One positive aspect about this approach is that it does not need to be developed for a specific dataset. Other approaches need to prepare statistical models (e.g. Bayesian networks) to model the real dataset, which then use it to generate the synthetic dataset, making it dependant on the correctness of this analysis. A GAN-based RSDGM does not depend on such models and will train its own (neural network) model, which can also be maintained as the real dataset changes. This allows to have a core approach that can be applied to multiple datasets without significant analysis. This was confirmed during the experiments, during which the GAN successfully generated realistic synthetic datasets based on three different datasets. An additional advantage of this is that the approach maintains associations in the data that may not be yet understood, and hence modelled manually. This can be particularly helpful for datasets used for discovery of new clinical knowledge. One identified challenge of the approach is the need for significant computing power to train the network and generate the dataset. However, this is not considered prohibitive. Future work of the justification includes further developing the arguments, also identifying specific evidence for the RSD based on a proof of concept. Future work will need to focus on further validation tests of the two datasets, benchmarking against other ML approaches and implementations, as well as systematically testing various configurations and hyper-parameters. Different architectures of the GAN can be explored by implementing other machine learning techniques such as SVMs for the discriminator, or auto-encoders for the generator. Finally, the work produced data in a format easily transformable to that of the input data; however larger scale adoption of the approach will need to be accompanied by more generalizable and accessible approach to packaging and accessing the data, which has been identified as a broader scope step. Overall, the GAN based RSDGM showed much promise and is considered a viable approach to be used for development of a healthcare dataset.

## Acknowledgements

This is a research product and does not reflect future policies nor recommendations of the involved organisations. Special thanks to Dr Eda Bilici Ozyigit for her technical work, as a research associate on the project.

## Declaration of conflicting interests

The author(s) declared no potential conflicts of interest with respect to the research, authorship, and/or publication of this article.

## Funding

The author(s) disclosed receipt of the following financial support for the research: This work was led by the MHRA and performed in collaboration with the NHS Digital and a2-ci under the £10m Regulators' Pioneer Fund, launched by The Department for Business, Energy and Industrial Strategy (BEIS). The fund enables UK regulators to develop innovation-enabling approaches to emerging technologies and unlock the long-term economic opportunities identified in the government's modern Industrial Strategy. Parts of this work were performed as part of Health Data Research (HDR) UK, which is funded by the UK Medical Research Council, Engineering and Physical Sciences Research Council, Economic and Social Research Council, Department of Health and Social Care (England), Chief Scientist Office of the Scottish Government Health and Social Care Directorates, Health and Social Care Research and Development Division (Welsh Government), Public Health Agency (Northern Ireland), British Heart Foundation and Wellcome Trust.

## Data statement

The results datasets for experiments 3-6 have been deposited at <http://wrap.warwick.ac.uk/162871/>. Access after approval by the authors due to their commercial in confidence nature.

## ORCID iDs

Theodoros N Arvanitis  <https://orcid.org/0000-0001-5473-135X>

George Despotou  <https://orcid.org/0000-0003-3437-6412>

## References

1. MHRA. Guidance: medical device stand-alone software including apps (including IVDMDs). Available: [https://assets.publishing.service.gov.uk/government/uploads/system/uploads/attachment\\_data/file/717865/Software\\_flow\\_chart\\_Ed\\_1-05.pdf](https://assets.publishing.service.gov.uk/government/uploads/system/uploads/attachment_data/file/717865/Software_flow_chart_Ed_1-05.pdf) (10/10/2019)
2. Regulation (EU) 2017/745 of the European Parliament and of the Council of 5 April 2017 on medical devices, amending Directive 2001/83/EC, Regulation (EC) No 178/2002 and Regulation (EC) No 1223/2009 and repealing Council Directives 90/385/EEC and 93/42/EEC.
3. FDA. Clinical Decision Support Software Draft Guidance for Industry and Food and Drug Administration Staff. Available: <https://www.fda.gov/media/109618/download> (10/10/2019).
4. Bellovin SM, Dutta PK and Reitering N. Privacy and synthetic datasets. *Stan Tech L Rev* 2019; 22: 1.
5. Buczak AL, Babin S and Moniz L. Data-driven approach for creating synthetic electronic medical records. *BMC Med Inform Decis Mak* 2010; 10: 59. DOI: [10.1186/1472-6947-10-59](https://doi.org/10.1186/1472-6947-10-59)
6. Moniz L, Buczak AL, Hung L, et al. Construction and validation of synthetic electronic medical records. *Online J Public Health Inform* 2009; 1:e2. DOI: [10.5210/ojphi.v1i1.2720](https://doi.org/10.5210/ojphi.v1i1.2720)
7. Walonoski J, Kramer M, Nichols J, et al. An approach, method, and software mechanism for generating synthetic patients and the synthetic electronic health care record, *J Am Med Inform Assoc* 2018; 25: 230–238. DOI: [10.1093/jamia/ocx079](https://doi.org/10.1093/jamia/ocx079)
8. Buczak AL, Babin S and Moniz L. Data-driven approach for creating synthetic electronic medical records, *BMC Med Inform Decis Making* 2010; 10: 59.
9. Thomson DR, Kools L and Jochem WC. Linking synthetic populations to household geolocations: A demonstration in Namibia. *Data* 2018; 3: 30.
10. Begoli E, Brown K, Srinivas S, et al. A generator framework for high-volume, high-fidelity synthetic mental health notes. 2018 IEEE International Conference on Big Data (Big Data), Seattle, WA, USA. 2018: 951–958.

11. Lee SH. Natural language generation for electronic health records. *npj Digital Med* 2018; 1: 63. DOI: [10.1038/s41746-018-0070-0](https://doi.org/10.1038/s41746-018-0070-0)
12. McLachlan S, Dube K and Gallagher T. Using the CareMap with health incidents statistics for generating the realistic synthetic electronic healthcare record. 2016 IEEE International Conference on Healthcare Informatics (ICHI), Chicago, IL. 2016: 439–448.
13. Dahmen J and Cook D. SynSys: a synthetic data generation system for healthcare applications. *Sensors* 2019; 19: 1181.
14. Avino L, Ruffini M and Gavalda R. *Generating synthetic but plausible healthcare record datasets*. arXiv: 1807.01514v1 [stat.ML] 4 Jul 2018
15. Wang Z, Myles P and Tucker A. Generating and evaluating synthetic UK primary care data: preserving data utility & patient privacy. 2019 IEEE 32nd International Symposium on Computer-Based Medical Systems (CBMS), Cordoba, Spain. 2019: 126–131. DOI: [10.1109/CBMS.2019.00036](https://doi.org/10.1109/CBMS.2019.00036).
16. Baoqaly MK, Liu C and Chen K. Realistic data synthesis using enhanced generative adversarial networks. 2019 IEEE Second International Conference on Artificial Intelligence and Knowledge Engineering (AIKE).
17. McLachlan S, Dube K, Gallagher T, et al. The ATEN framework for creating the realistic synthetic electronic health record. Proceedings of the 11th International Joint Conference on Biomedical Engineering Systems and Technologies, HEALTHINF, 5: 220–230.
18. Schiff S, Gerhke M and Moller R. Efficient enriching of synthesized relational patient data with time series data. *Proced Computer Sci* 2018; 141: 531–538.
19. Baowaly MK, Liu C and Chen K. Realistic data synthesis using enhanced generative adversarial networks. 2019 IEEE Second International Conference on Artificial Intelligence and Knowledge Engineering (AIKE), Sardinia, Italy, 2019: 289–292.
20. Norgaard S, Saeedi R, Sasani K, et al. Synthetic sensor data generation for health applications: a supervised deep learning approach. 2018 40th Annual International Conference of the IEEE Engineering in Medicine and Biology Society (EMBC), Honolulu, HI, 2018: 1164–1167.
21. Guan J, Li R, Yu S, et al. *Generation of synthetic electronic medical record text*. arXiv:1812.02793 [cs.CL]
22. Choi E, Biswal S, Malin B, et al. *Generating multi-label discrete patient records using generative adversarial networks*. arXiv:1703.06490 [cs.LG]
23. Shmelkov K, Schmid C and Alahari K. How good is my GAN? In: Proceedings Eur Conf Computer Vis (ECCV), 2018: 213–229.
24. Zare M and Wojtusiak J. Weighted Itemsets Error (WIE) approach for evaluating generated synthetic patient data. 2018 17th IEEE International Conference on Machine Learning and Applications (ICMLA), Orlando, FL, 2018: 1017–1022.
25. Borji A. Pros and cons of gan evaluation measures. *Computer Vis Image Understanding*. 2019; 179: 41–65.
26. NHS Digital. *Hospital Admitted Patient Care Activity, 2017-18, Publication, Part of Hospital Admitted Patient Care Activity, National statistics*. Publication 20 Sep 2018, Available at: <https://digital.nhs.uk/data-and-information/publications/statistical/hospital-admitted-patient-care-activity/2017-18>
27. NHS Digital. *Hospital Accident and Emergency Activity, 2017-18, Publication, Part of Hospital Accident & Emergency Activity*. Publication 13 Sep 2018, Available at: <https://digital.nhs.uk/data-and-information/publications/statistical/hospital-accident-emergency-activity/2017-18>
28. Johnson AEW, Pollard TJ, Shen L, et al. MIMIC-III, a freely accessible critical care database. *Scientific Data* 2016; 3: 160035. DOI: [10.1038/sdata.2016.35](https://doi.org/10.1038/sdata.2016.35).
29. Goodfellow I, Pouget-Abadie J, Mirza M, et al. Generative adversarial nets. In *Advances in neural information processing systems*. 2014, pp. 2672–2680.



30. Mirza M and Osindero S. Conditional generative adversarial nets. arXiv preprint arXiv:1411.1784. 2014 Nov 6.
31. Arjovsky M, Chintala S and Bottou L. Wasserstein gan. arXiv preprint arXiv:1701.07875. 2017 Jan 26.
32. Gulrajani I, Ahmed F, Arjovsky M, et al. Improved training of wasserstein gans. *In Advances in neural information processing systems*. 2017; 5767–5777.
33. Laleci Erturkmen GB, Yuksel M, Sarigul B, et al. A collaborative platform for management of chronic diseases via guideline-driven individualized care plans. *Comput Struct Biotechnol J* 2019; 17: 869–885. DOI:[10.1016/j.csbj.2019.06.003](https://doi.org/10.1016/j.csbj.2019.06.003)
34. Kingma DP and Ba J. Adam: A method for stochastic optimization. arXiv preprint arXiv:1412.6980. 2014 Dec 22.
35. Ramachandran P, Zoph B and Le QV. *Searching for activation functions*, <https://arxiv.org/abs/1710.05941v2>, 16 Oct 2017.
36. Despotou G, Harrison S, White S, et al. Safety justification of healthcare applications using synthetic datasets. In: *The Importance of Health Informatics in Public Health during a Pandemic. Studies in Health Technology and Informatics*, IOS Press, 272, pp. 35–38.